

# Recommender Systems - using Linear Algebra

Giri Iyengar

# Introductions

- Chief Architect at AOL Platforms. Machine Learning startup recently acquired by AOL/Verizon
- Over 2 decades of building models and teaching computers to do interesting things
- Electrical Engineer, transitioned to Pattern Recognition/Machine Learning during my doctoral studies
- IIT Mumbai, MIT, IBM Research, 2 start-ups
- Designed & Teaching this course for the past 4 terms

# Course Outline

## Linear Algebra

- Matrices and Vectors
- Some standard Linear Algebra techniques
- Applications to Data Analytics & Machine Learning

## Probability and Statistics

- Probability Theory, Independence, Conditional Independence, Bayes Rule, etc
- Statistics, Hypothesis testing, Statistical Learning
- Frequentist vs Bayesian points of view

## Calculus and Numerical Techniques

- Quick review of Calculus and Numerical techniques
- Gradient Descent, L-BFGS for machine learning

# Assignments and Grading

- 1 Assignment every week
- 2 Problem sets. First is quiz-style and the second is a bit more involved – will ask you to code
- 1 Final exam. Take home style. You'll have a weekend to work on it
- Encourage you to share, discuss, and build a community

# Recommender Systems

Let's set up a toy recommendation problem. There are 4 users, reading 5 different blogs. Here is a data frame that contains information about how many times each person read these 5 blogs.

`http://www.ibm.com/developerworks/library/os-recommender1/index.html`

# Sample dataset

<b>Blogs</b>	Marc	Megan	Elise	Jill
Linux	13	3	11	-
Open Source	10	-	-	3
Cloud Computing	6	1	9	-
Java	-	6	-	9
Agile	-	7	1	8

# Reading habits of Users

```
Marc <- c(13,10,6,0,0)
Megan <- c(3,0,1,6,7)
Elise <- c(11,0,9,0,1)
Jill <- c(0,3,0,9,8)
u <- data.frame(Marc,Megan,Elise,Jill)
rownames(u) <- c('Linux','Open source','Cloud computing','Java',
u
```

##	Marc	Megan	Elise	Jill
## Linux	13	3	11	0
## Open source	10	0	0	3
## Cloud computing	6	1	9	0
## Java	0	6	0	9
## Agile	0	7	1	8

# User-Item matrix


Construct a User-Item matrix where each row is a user. Raw counts are not so useful, so normalize them as well.

```
unorm <- u
for (i in 1:length(u)) { unorm[,i] <- u[,i] / sqrt(u[,i] %*% u[,i]) }
t(unorm)
```


	Linux	Open source	Cloud computing	Java
Marc	0.7443778	0.5725983	0.3435590	0.0000000
Megan	0.3077935	0.0000000	0.1025978	0.6155870
Elise	0.7720486	0.0000000	0.6316762	0.0000000
Jill	0.0000000	0.2417469	0.0000000	0.7252407



# Measuring Proximity between Users



# Measuring Proximity between Users



# Measuring Proximity between Users

- Cosine of the angle between these vectors can be used. Also known as cosine distance.
- The cosine distance between users can be elegantly expressed as  $unorm' \times unorm$

```
as.matrix(t(unorm)) %*% as.matrix(unorm)
```

```
##           Marc      Megan      Elise      Jill
## Marc  1.0000000  0.2643631  0.79171393  0.13842387
## Megan 0.2643631  1.0000000  0.35284686  0.90943261
## Elise 0.7917139  0.3528469  1.00000000  0.04524615
## Jill  0.1384239  0.9094326  0.04524615  1.00000000
```

- Comparing cosine distances, we see that it agrees with the scatterplot of the reading habits